

Hippo

finding business ideas using AI and big data analysis

Architecture Document

Document revision 2.0

13 March 2018

Client

Juicy Story

Teaching Assistant

Hichem Bouakaz

Team

Levi van Rheenen

Jean Paul Donovan Meijer

Said Faroghi

Natalia Karpova

Thijs Baksteen

Andreea Glavan

Sardor Khashimov

Table of contents

Table of contents	2
Introduction	3
Architectural Overview	4
Technology Stack	6
Data analysis	6
Front-end development	7
Back-end development	7
Team Organisation	9
Change Log	10

Introduction

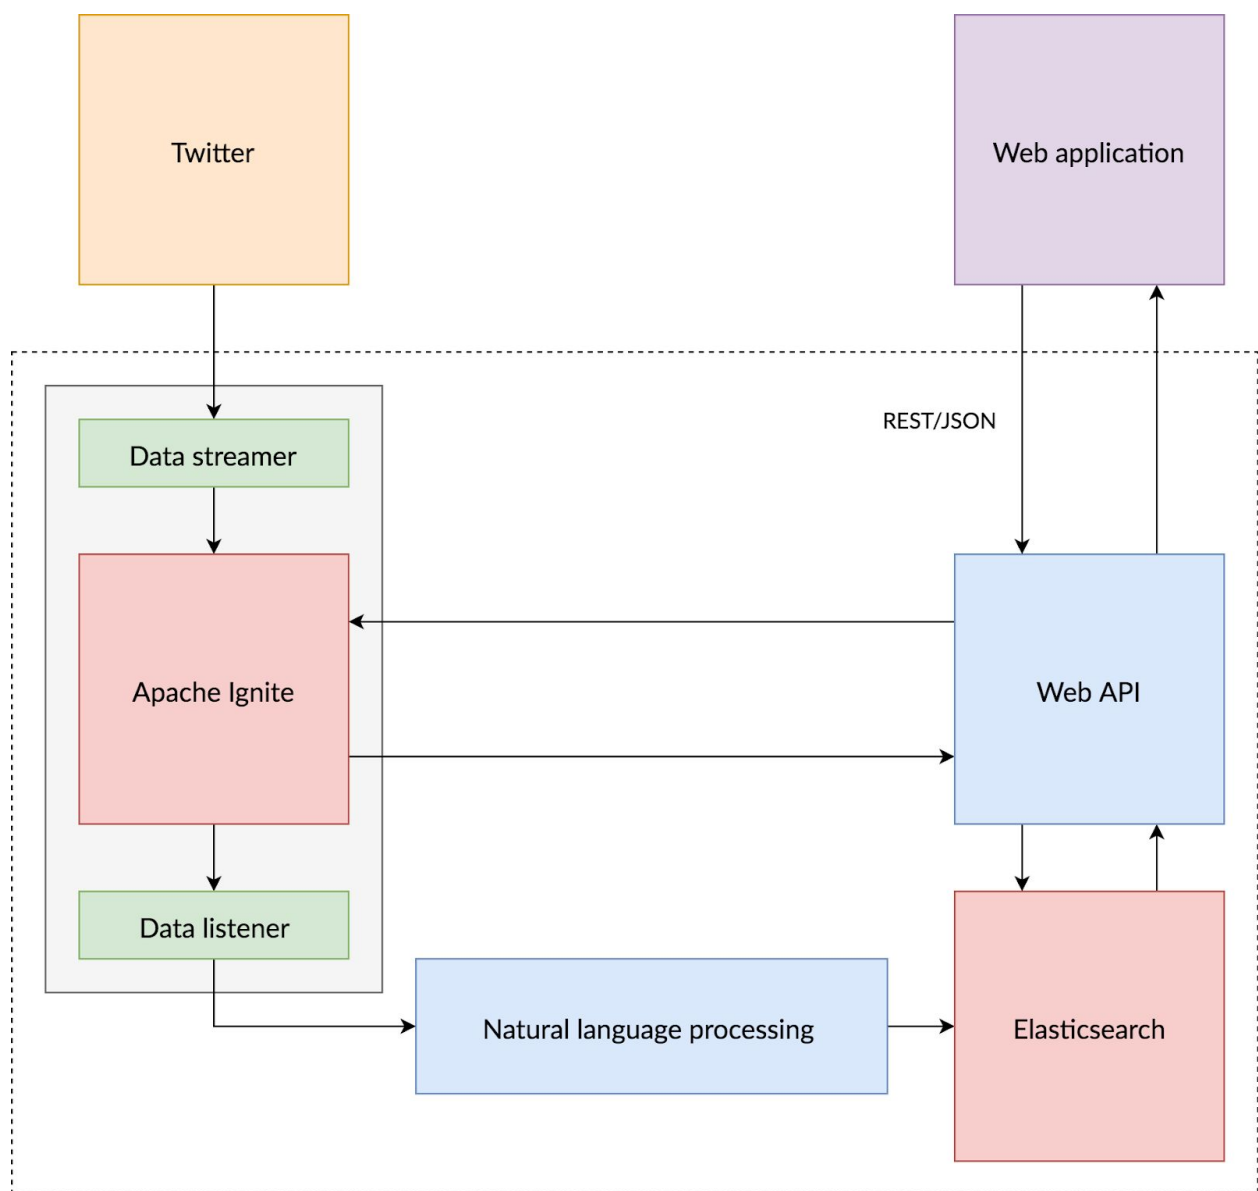
Hippo is a web based application for entrepreneurs. Entrepreneurs can search for keywords and find tweets that display desire that have to do with those keywords. In order to do this, we need to get data from twitter, filter it, be able to search through it, and display it on a website.

We split the project into three separate parts: front end, data processing, and back end. The front end will focus on the interface of the website and user interaction. It will get its data from the back end, which is going to focus on the data delivery. The data processing part will input data from Twitter, filter it, and process it. The back end will search through the data gathered from this processing part.

This document will delve deeper into these parts. It will go into the implementation and design details, and also give some insight on the decisions made.

Architectural Overview

The application consists out of three major components, that are data analysis, the web API and the web application. Data analysis can be divided into three smaller components which all have their own function. The architecture of the other two main components is more trivial, and are both eventually just single applications. Furthermore, both data analysis and the web API will make use of the database systems Apache Ignite and Elasticsearch.



The architecture of data analysis can be split into three components, which are actually all separate services. The data streamer connects to the Twitter streaming API and stores the incoming data stream in Apache Ignite, the database we are using. The data listener sends the new data added to Ignite to natural language processing over HTTP. This then processes the tweets, using natural language processing techniques and stores the all relevant data in Elasticsearch.

The web API provides functionality to search and present the data that has been analysed. This is initiated by calls from the web application which runs in the users internet browser. The web application provides an interface the user can use to interact with any data stored on the platform. The web API and the web application communicate over REST where content will be encode in JSON format.

Connections between the databases and the components will be done using the official client libraries of Apache Ignite and Elasticsearch.

Technology Stack

Data analysis

The data will be captured in real time from Twitter, the (real time) streaming API is less restricted by rate limiting. Practically eliminating this burden and allowing us to access a much greater variety of data to search through. This data is captured by a small application written in Java that uses the data streamer for Twitter that is already provided by Apache Ignite. The raw data from Twitter is will be stored in the database, and processed afterwards when possible.

Apache Ignite allows us easily handle a growing data set in the future as well allow further computations on them later. Any changes in the database are reported to the data listener, a small application in Java, which then pushes the new entity to data analysis, which performs natural language processing on the tweets contents, this is a Python application.

For data analysis part Python programming language together with NLTK (Natural Language Toolkit) platform will be used. Natural language processing will be done in several steps. This is a preliminary layout of an analysation process and may change during the later development and optimization of Hippo application.

Both users search queueries and tweets themselves will be analyzed using nltk. As the first step, keywords extraction will be performed on user queueries. This will include several steps, namely, tokenization (the exact tokenizer will be determined later by experimenting with the data) of a queuery, POS-tagging (part-of-speech tagging), named entity detection (using, presumably, chunking, tag patterns and chinking) and relation detection (at the early stages it would be done with rule-based system). WordNET will be used to identify the exact meaning of keywords and to find synonyms to improve consequent search.

The same keywords extraction will be performed on tweets. In case of tweet analysis, relation detection part will be subjective to particular search query entered by a user. We also consider to make tweets filtration based on their similarity to each other so that user will get a bunch of different ideas as an output.

Pre-fatorial strategy is following:

Search query by itself may act as centroid document. Each tweets that passes previous “filter” , in other words, each tweet which keywords are related to search query keywords, will be treated as a separate “document” and will be compared to centroid document and other “documents”

(tweets). This will be done with a version of CSIS (Cross Sentence Information Subsumption) and will result in exclusion of identical ideas.

Front-end development

The front-end side of the project will mostly use Typescript (in conjunction with HTML & CSS) to build the web app. Typescript, being a superset of Javascript, adds a lot of extra functionality such as the ability to easily code in an Object Oriented manner. This includes concepts like classes and interfaces. Moreover, Typescript has a better compiler that facilitates debugging. (Note on early iterations: Because the front end will be relatively simple, production may start using Javascript at first and then converted to typescript in the later stages). The choice of HTML and CSS is pretty clear, since they are standard in web development.

We will use two Frameworks - Vue.js with webpack for JavaScript and Foundation for CSS. Vue.js is currently one of the industry's favorite frameworks because it has clear advantages over others (Angular, React), and is simple to learn and code in. Vue.js will also be coupled with the third-party Axios, a promise-based HTTP client library for REST api consumption. Vue had "vue-resource" that was able to do this, but since Vue 2.0 this has been retired, and the Vue.js + Axios combo is the most recommended setup.

Foundation was chosen over other frameworks like Bootstrap. While Bootstrap may allow faster prototyping with many templates, Foundation offers a more "do it yourself" approach that can offer a unique experience to the user and also allow effective customization for the developer. Thus, as customizable features are valued, Foundation is an ideal choice.

Front end unit testing will be done with Karma.

Back-end development

For our back-end, we will use Python with Flask. Python is chosen because most of the team is familiar with the language, and it serves our purposes well. Flask is a microframework for Python, meaning that it provides a core framework that is relatively small, but extensible with a variety of different libraries. This avoids the clutter one would have with a larger framework that includes functionality that is not needed for the application.

Elasticsearch will be used for our user searching needs. Elasticsearch is an open source search engine/database that can store the many ideas we collect, and allow search options to find ideas that the user wants.

We will use Docker to support our architecture, and provide secure containers for the different components of the back-end to run in. Docker will also simplify the process of scaling the application to more servers if needed in the future.

Team Organisation

We chose to split into 2 teams: front end and back end, with 2 and 5 people, respectively. The reason for not choosing a 3:4 split was due to the fact that the front end design is on the minimalistic side, based on the demos we have seen.

Our front end team consists of Said Faroghi and Sardor Khashimov, and our back end team consists of Levi van Rheenen, Jean P.D. Meijer, Natalia Karpova, Thijs Baksteen, and Andreea Glavan.

As the names indicate, the first team will be in charge of the front end which includes the web design, implementation of this design, creating a logo, and maintaining the site.

The back end team is in charge of data processing and web back end. The back end focus is on data processing and filtering of ideas and categories, as we believe this is the key to achieving the end goals of this project.

The data processing part includes creating a data streamer to crawl Twitter for ideas, updating the database with the received ideas, calculate each ideas' stat points, filtering these ideas into their respective categories, and filtering ideas which are the same but worded differently into the same database entry.

The web back end is in charge of communicating with the front end and sending ideas from the database based on the given key words. What is more, the web back end deals with users' accounts and stats functionality.

Change Log

The record of the changes made to the architecture document tagged with date and author.

Author	Date	Description
Jean P.D. Meijer	10-03-2018	Create the architecture document and layout the structure.
Andreea Glavan	10-03-2018	Updated the team organisation section.
Natalia	11-03-2018	Added section about the data analysis tech stack.
Said & Sardor	12-03-2018	Added description for front-end tech stack.
Levi	13-03-2018	Added introduction.
Andreea Glavan	13-03-2018	Reviewed and updated existing sections.
Jean P.D. Meijer	13-03-2018	Added the architecture overview section and extended the data analysis tech stack.
Thijs	13-03-2018	Added description for back-end tech stack.